

1. The knowledge acquired during training is embedded in the model parameters and used at inference time, rather than retrieved from the original dataset.
2. Identical inputs do not always lead to identical outputs when interacting with a language model.
3. The behaviour of a language model may be influenced by provider-defined instructions that are not visible to the user.
4. Outputs produced by language models can reflect patterns and biases present in the data they were trained on.
5. Language models generate text incrementally by selecting likely continuations given the context.
6. The use of language models does not remove the researcher's responsibility for the validity of the results.
7. Increasing model size typically comes with higher computational and infrastructure demands.